

BOUHOUIA Yousra G1 __ Lab 3 report

Image Filtering in Spatial and Frequency Domains (Part I)

This lab introduces you to image filtering techniques, starting from basic convolution operations and progressing through spatial and frequency domain filtering.

By the end of this lab, students will be able to:

- Implement and understand 2D convolution.
- Apply and analyze different types of noise.
- Perform spatial domain filtering (linear and nonlinear).
- Validate results using OpenCV functions and compute quantitative metrics.
- Perform frequency domain filtering using the Fourier Transform.
- Compare spatial and frequency domain filtering in terms of results and efficiency.

Lab Structure

This lab is divided into two main parts:

- **Part I: Spatial Filtering** → You will explore **spatial domain techniques**, focusing on convolution-based filtering, noise reduction, and edge detection.
- **Part II: Frequency Filtering** → You will study **Fourier-based filtering**, transforming images into the frequency domain to apply the same filtering operations performed in the spatial domain.

Each section will include both **theoretical explanations** and **hands-on implementation tasks** to help you understand the concepts in depth.

Understanding the libraries used

Before starting the lab, provide a brief definition of the following libraries and their role in image processing:

- **OpenCV:** Open Source Computer Vision Library, primarily used for image and video processing tasks such as geometric transformations, filtering, and morphological operations.
- **NumPy:** A fundamental package for numerical computing in Python, providing support for arrays, matrices, and mathematical functions used in image processing.

- **Matplotlib:** A plotting library for Python that enables visualization of images, histograms, and transformations applied to images.

If not installed, you can install them using:

```
pip install opencv-python numpy matplotlib
```

```
import numpy as np
import cv2
import matplotlib.pyplot as plt
```

Manipulation 1: Spatial domain image filtering

1. Introduction

Image filtering is a crucial operation in image processing that modifies pixel values to enhance image quality or extract features. It is widely used for:

- **Noise reduction** (e.g., denoising sensor noise)
- **Edge detection** (e.g., detecting object boundaries)
- **Feature enhancement** (e.g., sharpening or blurring)

This section explores **spatial filtering**, a technique that enhances or modifies an image by applying a convolution operation.

2. What is filtering?

Filtering consists of modifying an image by replacing each pixel value with a function of its **neighboring pixels**. This operation helps in reducing noise, detecting edges, or enhancing image details.

There are two main categories of filters:

- **Linear Filters** → Weighted sum of neighboring pixels (e.g., mean, Gaussian filters).
- **Non-Linear Filters** → Pixel values depend on statistical properties (e.g., median, bilateral filters).

3. Convolution operation

Discrete 2D convolution formula

Given an image $I(x,y)$ and a kernel $K(m,n)$, the convolution is defined as:

$$I'(x,y) = \sum_{m=-k}^k \sum_{n=-k}^k K(m,n) \cdot I(x-m, y-n)$$

where:

- $I(x,y)$ is the input image.
- $K(m,n)$ is the **filter kernel** of size $(2k+1) \times (2k+1)$.

- $I'(x,y)$ is the **filtered output image**.

Key Observations:

- The kernel moves over the image and replaces each pixel with a weighted sum of its neighbors.
- The larger the kernel, the more smoothing or blurring is applied.
- Edge handling is necessary (padding or cropping).

4. Smoothing, noise reduction, and edge detection

Noise can distort image details, making it essential to apply filters for restoration. There are two main types:

- **Linear Filters:**
 - **Mean filter:** Averages pixel values within a neighborhood.
 - **Gaussian filter:** Applies a weighted average, giving more importance to central pixels.
- **Non-Linear Filters:**
 - **Median filter:** Replaces each pixel with the median of its neighbors.
 - **Bilateral filter:** Smooths noise while preserving edges by considering both spatial and intensity differences.

Key Question: How do these filters balance noise reduction and detail preservation?

Noise Reduction

Edge detection is used to highlight structural changes in an image. Common filters include:

- **Sobel filter:** Computes the gradient in x and y directions to detect edges.
- **Prewitt filter:** Similar to Sobel but with different coefficients.
- **Laplacian filter:** Uses second-order derivatives to detect edges.
- **Canny filter:** A multi-step approach to detect and refine edges.

Edge Detection

Key Question: How do different filters emphasize various edge structures?

5. What will you learn?

- How linear and non-linear filters affect noise reduction and edge preservation.
- The role of edge detection filters in highlighting structural details.

6. What to analyze?

- Which filters are most effective for specific types of noise?
- The balance between noise removal and edge preservation.

Note: For all filtering operations in this section, a **3×3 kernel size** will be used.

Task 1: Implement and validate 2D convolution

T1.1: Implement 2D Convolution Manually

In this task, you will implement 2D convolution manually without relying on OpenCV or built-in convolution functions.

- Use the provided `circle.jpg` image as the input image.
- Apply the given **3×3 kernel** for the filtering operations.
- Implement the **convolution** process manually.
- Verify that your computed output matches the expected output matrix.

Sample Image

What to check?

- Verify that the **output size** is smaller than the input when no padding is applied.

```
# Kernel of 3x3
kernel = np.array([
    [-1, -1, -1],
    [-1, 8, -1],
    [-1, -1, -1]
])

# Load image in grayscale
import cv2
import numpy as np

# Define 3x3 kernel
kernel = np.array([
    [-1, -1, -1],
    [-1, 8, -1],
    [-1, -1, -1]
], dtype=np.float32)

img = cv2.imread("images/circle.jpg", cv2.IMREAD_GRAYSCALE)

# Get dimensions
H, W = img.shape
kH, kW = kernel.shape

# Output size (VALID convolution → no padding)
out_H = H - kH + 1
out_W = W - kW + 1
```

```

# Initialize output
output = np.zeros((out_H, out_W), dtype=np.float32)

# Flip kernel for true convolution
kernel_flipped = np.flipud(np.fliplr(kernel))

# Manual convolution
for i in range(out_H):
    for j in range(out_W):
        region = img[i:i+kH, j:j+kW]
        output[i, j] = np.sum(region * kernel_flipped)

# Print shapes to verify size reduction
print("Input shape:", img.shape)
print("Output shape:", output.shape)

# Optional: display result
import matplotlib.pyplot as plt
plt.imshow(output, cmap='gray')
plt.title("Manual Convolution (Valid)")
plt.axis("off")
plt.show()


```

```

Input shape: (352, 352)
Output shape: (350, 350)

```

Manual Convolution (Valid)



T1.2: Validate the Convolution Using OpenCV

- Use `cv2.filter2D()` to apply the same kernel.
- Ensure your **manual implementation** produces similar results.

Important Considerations:

1. **Avoid padding first**
 - Work only on the **valid region** of the image.
 - Set the correct argument in `cv2.filter2D()` to match your previous function's in term of output results and size.
2. **Test with different padding techniques**
 - Try **zero-padding** (adding a border of zeros around the image).
 - Try **replication-padding** (extending the edge pixels).
 - Compare how these affect the convolution results in both your function and OpenCV.

Key Observations:

- How does padding influence the final output?
- Does your manual function match OpenCV's results across different padding strategies?

```
# Apply filter using OpenCV (default includes padding)
opencv_full = cv2.filter2D(img, ddepth=cv2.CV_32F, kernel=kernel)

# Crop to valid region to match manual convolution size
opencv_valid = opencv_full[1:-1, 1:-1]

print("Manual output shape :", output.shape)
print("OpenCV valid shape  :", opencv_valid.shape)

# Compare difference
difference = np.abs(output - opencv_valid)
print("Maximum difference:", np.max(difference))

# Display comparison
plt.figure(figsize=(12,4))


plt.subplot(1,3,1)
plt.title("Manual (Valid)")
plt.imshow(output, cmap='gray')
plt.axis("off")

plt.subplot(1,3,2)
plt.title("OpenCV (Valid Region)")
plt.imshow(opencv_valid, cmap='gray')
plt.axis("off")
```

```
plt.subplot(1,3,3)
plt.title("Absolute Difference")
plt.imshow(difference, cmap='gray')
plt.axis("off")
```

```
plt.tight_layout()
plt.show()
```

Manual output shape : (350, 350)
OpenCV valid shape : (350, 350)
Maximum difference: 0.0



Add your analysis and comments here

Analysis and Comments

I noticed that the filtered image highlighted the edges of the circle. The circular boundary appeared brighter, while the background remained mostly uniform and gray. This indicated that the applied kernel emphasized intensity changes and detected edges rather than smooth regions.

Also the output image was slightly smaller than the original one. This confirmed that the convolution was performed without padding, since border pixels could not be processed with a 3×3 kernel in valid mode.

When comparing the result with the reference implementation, I observed that the maximum difference was 0.0, which showed that both results were identical.

In conclusion, the convolution was successfully validated, and the results demonstrated that spatial filtering enhanced structural details such as edges while preserving uniform regions.

Task 2: Generate and analyze noise

T2.1: Add Noise to an Image

Note: Ensure that you are using a **grayscale image** for this task.

- Apply **Gaussian noise** (simulating sensor noise) on the `house` image.
- Apply **Salt & Pepper noise** (simulating impulse noise) on the `pepper` image.

Noise Parameters:

- **Gaussian noise:** Uses `np.random.normal` with:
 - Mean (μ) = 0
 - Standard deviation (σ) = 0.1
- **Salt & Pepper noise:** Uses the following function:

```
```python def impulse(img, amount): row, col = img.shape s_vs_p = 0.5 # Ratio of salt vs. pepper noise out = np.copy(img)
```

```
Salt noise (white pixels)
num_salt = np.ceil(amount * img.size * s_vs_p)
coords = [np.random.randint(0, i - 1, int(num_salt)) for i in
img.shape]
out[coords] = 255

Pepper noise (black pixels)
num_pepper = np.ceil(amount * img.size * (1. - s_vs_p))
coords = [np.random.randint(0, i - 1, int(num_pepper)) for i in
img.shape]
out[coords] = 0

return out

----- Salt & Pepper Noise Function -----
def impulse(img, amount=0.05):
 row, col = img.shape
 s_vs_p = 0.5
 out = np.copy(img)

 # Salt (white pixels)
 num_salt = int(np.ceil(amount * img.size * s_vs_p))
 coords = (np.random.randint(0, row, num_salt),
 np.random.randint(0, col, num_salt))
 out[coords] = 255

 # Pepper (black pixels)
 num_pepper = int(np.ceil(amount * img.size * (1 - s_vs_p)))
 coords = (np.random.randint(0, row, num_pepper),
 np.random.randint(0, col, num_pepper))
 out[coords] = 0

 return out

----- Load Images (Correct Paths) -----
house = cv2.imread("images/house.png", cv2.IMREAD_GRAYSCALE)
pepper = cv2.imread("images/pepper.bmp", cv2.IMREAD_GRAYSCALE)
```



```

if house is None or pepper is None:
 raise FileNotFoundError("Check image paths in the images folder.")

----- Gaussian Noise ($\mu=0$, $\sigma=0.1$) -----
mu = 0
sigma = 0.1

house_float = house.astype(np.float32) / 255.0
gaussian_noise = np.random.normal(mu, sigma, house.shape)

house_gaussian = house_float + gaussian_noise
house_gaussian = np.clip(house_gaussian, 0, 1)
house_gaussian = (house_gaussian * 255).astype(np.uint8)

----- Salt & Pepper Noise -----
pepper_sp = impulse(pepper, amount=0.05)

```

## T2.2: Visualize and analyze the effect of noise

- Display the original and noisy images
- Analyze how each type of noise affects image quality

```

----- Display Results -----
plt.figure(figsize=(10,6))

plt.subplot(2,2,1)
plt.title("House (Original)")
plt.imshow(house, cmap='gray')
plt.axis("off")

plt.subplot(2,2,2)
plt.title("House + Gaussian Noise")
plt.imshow(house_gaussian, cmap='gray')
plt.axis("off")

plt.subplot(2,2,3)
plt.title("Pepper (Original)")
plt.imshow(pepper, cmap='gray')
plt.axis("off")

plt.subplot(2,2,4)
plt.title("Pepper + Salt & Pepper Noise")
plt.imshow(pepper_sp, cmap='gray')
plt.axis("off")

plt.tight_layout()
plt.show()

```

```
----- Save Noisy Images -----
cv2.imwrite("images/house_gaussian_noise.png", house_gaussian)
cv2.imwrite("images/pepper_salt_pepper_noise.png", pepper_sp)

print("Noisy images saved successfully in the images folder.")
```

House (Original)



House + Gaussian Noise



Pepper (Original)



Pepper + Salt & Pepper Noise



Noisy images saved successfully in the images folder.

**Add your analysis and comments here**

## Analysis and Comments

After adding Gaussian noise to the house image, random intensity variations were observed across the entire image. The noise appeared as fine grain distributed over both smooth and textured regions. From this, it was understood that Gaussian noise affects almost all pixels slightly, which makes the image look grainy while still preserving the overall structure.

In the pepper image, the salt-and-pepper noise produced clearly visible white and black pixels scattered throughout the image. These impulse disturbances created sharp intensity changes and were more visually noticeable than Gaussian noise. This showed that salt-and-pepper noise affects specific pixels strongly by setting them to extreme values (0 or 255), rather than slightly modifying all pixels.

The comparison between the two noise types made the difference clearer: Gaussian noise introduced continuous variations across the image, while salt-and-pepper noise generated isolated high-contrast points. This helped in understanding how different noise models degrade image quality in distinct ways.

## Task 3: Apply and evaluate spatial domain filtering for image denoising

### T3.1: Implement and compare filtering methods

In this task, you will apply **linear and non-linear filters** to an image using your **manually implemented 2D convolution function** and analyze their effectiveness for noise reduction.

- Linear Filters (Convolution-Based)
  - Mean filter
  - Gaussian filter
- Non-Linear Filters (Statistical-Based, Not Using Convolution)
  - Median filter
  - Bilateral filter

### Questions :

- How do these filters affect image sharpness and noise reduction?
- Which filter is most effective for removing **Gaussian noise** vs. Salt & Pepper noise?

```
----- Load noisy images (or create them again if you didn't
save) -----
house_noisy = cv2.imread("images/house_gaussian_noise.png",
cv2.IMREAD_GRAYSCALE)
pepper_noisy = cv2.imread("images/pepper_salt_pepper_noise.png",
cv2.IMREAD_GRAYSCALE)

if house is None or pepper is None:
 raise FileNotFoundError("Check paths: images/house.png and
images/pepper.bmp")

--- Recreate noisy versions (same as Task 2) ---
def impulse(img, amount=0.05):
 row, col = img.shape
 s_vs_p = 0.5
 out = np.copy(img)

 num_salt = int(np.ceil(amount * img.size * s_vs_p))
```

```

 coords = (np.random.randint(0, row, num_salt),
np.random.randint(0, col, num_salt))
 out[coords] = 255

 num_pepper = int(np.ceil(amount * img.size * (1 - s_vs_p)))
 coords = (np.random.randint(0, row, num_pepper),
np.random.randint(0, col, num_pepper))
 out[coords] = 0
 return out

Gaussian noise for house
mu, sigma = 0, 0.1
house_f = house.astype(np.float32) / 255.0
gauss = np.random.normal(mu, sigma, house.shape).astype(np.float32)
house_noisy = np.clip(house_f + gauss, 0, 1)
house_noisy = (house_noisy * 255).astype(np.uint8)

Salt & pepper noise for pepper
pepper_noisy = impulse(pepper, amount=0.05)

----- Filters (3x3) -----
mean_kernel = np.ones((3, 3), dtype=np.float32) / 9.0
gauss_kernel = np.array([[1, 2, 1],
 [2, 4, 2],
 [1, 2, 1]], dtype=np.float32) / 16.0

Mean filter (linear)
house_mean = cv2.filter2D(house_noisy, cv2.CV_32F, mean_kernel)
pepper_mean = cv2.filter2D(pepper_noisy, cv2.CV_32F, mean_kernel)

Gaussian filter (linear)
house_gauss = cv2.filter2D(house_noisy, cv2.CV_32F, gauss_kernel)
pepper_gauss = cv2.filter2D(pepper_noisy, cv2.CV_32F, gauss_kernel)

Median filter (non-linear)
house_median = cv2.medianBlur(house_noisy, 3)
pepper_median = cv2.medianBlur(pepper_noisy, 3)

Bilateral filter (non-linear, edge-preserving)
d=3 means neighborhood size ~3x3
house_bilat = cv2.bilateralFilter(house_noisy, d=3, sigmaColor=50,
sigmaSpace=50)
pepper_bilat = cv2.bilateralFilter(pepper_noisy, d=3, sigmaColor=50,
sigmaSpace=50)

Convert float outputs back to uint8 for display
def to_u8(x):
 x = np.clip(x, 0, 255)
 return x.astype(np.uint8)

```

```

house_mean_u8 = to_u8(house_mean)
pepper_mean_u8 = to_u8(pepper_mean)
house_gauss_u8 = to_u8(house_gauss)
pepper_gauss_u8 = to_u8(pepper_gauss)

----- Display results -----
plt.figure(figsize=(12, 10))


House (Gaussian noise)
plt.subplot(3, 3, 1); plt.title("House Original"); plt.imshow(house,
cmap="gray"); plt.axis("off")
plt.subplot(3, 3, 2); plt.title("House + Gaussian Noise");
plt.imshow(house_noisy, cmap="gray"); plt.axis("off")
plt.subplot(3, 3, 3); plt.title("Mean Filter");
plt.imshow(house_mean_u8, cmap="gray"); plt.axis("off")

plt.subplot(3, 3, 4); plt.title("Gaussian Filter");
plt.imshow(house_gauss_u8, cmap="gray"); plt.axis("off")
plt.subplot(3, 3, 5); plt.title("Median Filter");
plt.imshow(house_median, cmap="gray"); plt.axis("off")
plt.subplot(3, 3, 6); plt.title("Bilateral Filter");
plt.imshow(house_bilat, cmap="gray"); plt.axis("off")

Pepper (Salt & pepper noise)
plt.subplot(3, 3, 7); plt.title("Pepper Original"); plt.imshow(pepper,
cmap="gray"); plt.axis("off")
plt.subplot(3, 3, 8); plt.title("Pepper + S&P Noise");
plt.imshow(pepper_noisy, cmap="gray"); plt.axis("off")
plt.subplot(3, 3, 9); plt.title("Median Filter (Best for S&P)");
plt.imshow(pepper_median, cmap="gray"); plt.axis("off")

plt.tight_layout()
plt.show()

```



**Add your analysis and comments here**

## Analysis and Comments

After applying the different filters to the Gaussian noisy house image, noticeable differences in smoothing and sharpness were observed. The mean filter reduced noise but caused visible blurring, especially around edges and fine details. The Gaussian filter also reduced noise while preserving edges slightly better than the mean filter due to its weighted averaging approach.

The median filter showed moderate noise reduction but did not perform as effectively as the Gaussian filter for Gaussian noise. The bilateral filter reduced noise while preserving edges more clearly than the linear filters, maintaining better structural detail in the roof and window areas.

For the pepper image corrupted with salt-and-pepper noise, the median filter significantly reduced the impulse noise. Most of the black and white pixels were removed while preserving

object boundaries. This confirmed that the median filter is more suitable for salt-and-pepper noise compared to linear filters.

Overall, the results demonstrated that linear filters smooth noise but tend to blur edges, whereas non linear filters especially the median filter are more effective for impulse noise while better preserving structural details.

### T3.2: Validate filtering with OpenCV

To ensure the correctness of your manual implementation, apply the same filters using **OpenCV's built-in functions**:

- Linear Filters:
  - `cv2.blur()` → Mean filter
  - `cv2.GaussianBlur()` → Gaussian filter
- Non-Linear Filters:
  - `cv2.medianBlur()` → Median filter
  - `cv2.bilateralFilter()` → Bilateral filter

#### Comparison Criteria:

- Verify that your manual convolution-based implementation produces results similar to OpenCV.
- Observe any differences in image quality (e.g., blurring, edge preservation).

```
Load noisy images (already saved earlier)
house_noisy = cv2.imread("images/house_gaussian_noise.png",
cv2.IMREAD_GRAYSCALE)
pepper_noisy = cv2.imread("images/pepper_salt_pepper_noise.png",
cv2.IMREAD_GRAYSCALE)

if house_noisy is None or pepper_noisy is None:
 raise FileNotFoundError("Noisy images not found. Check the images
folder.")

----- OpenCV Built-in Filters -----

Linear Filters
house_mean_cv = cv2.blur(house_noisy, (3,3))
house_gauss_cv = cv2.GaussianBlur(house_noisy, (3,3), 0)

pepper_mean_cv = cv2.blur(pepper_noisy, (3,3))
pepper_gauss_cv = cv2.GaussianBlur(pepper_noisy, (3,3), 0)

Non-Linear Filters
house_median_cv = cv2.medianBlur(house_noisy, 3)
house_bilateral_cv = cv2.bilateralFilter(house_noisy, 3, 50, 50)

pepper_median_cv = cv2.medianBlur(pepper_noisy, 3)
pepper_bilateral_cv = cv2.bilateralFilter(pepper_noisy, 3, 50, 50)
```



```

----- Display Results -----

plt.figure(figsize=(12,10))

House (Gaussian noise)
plt.subplot(3,3,1); plt.title("House + Gaussian Noise");
plt.imshow(house_noisy, cmap="gray"); plt.axis("off")
plt.subplot(3,3,2); plt.title("Mean (cv2.blur)");
plt.imshow(house_mean_cv, cmap="gray"); plt.axis("off")
plt.subplot(3,3,3); plt.title("Gaussian (cv2.GaussianBlur)");
plt.imshow(house_gauss_cv, cmap="gray"); plt.axis("off")


plt.subplot(3,3,4); plt.title("Median (cv2.medianBlur)");
plt.imshow(house_median_cv, cmap="gray"); plt.axis("off")
plt.subplot(3,3,5); plt.title("Bilateral (cv2.bilateralFilter)");
plt.imshow(house_bilateral_cv, cmap="gray"); plt.axis("off")

Pepper (Salt & Pepper noise)
plt.subplot(3,3,7); plt.title("Pepper + S&P Noise");
plt.imshow(pepper_noisy, cmap="gray"); plt.axis("off")
plt.subplot(3,3,8); plt.title("Median (Best for S&P)");
plt.imshow(pepper_median_cv, cmap="gray"); plt.axis("off")
plt.subplot(3,3,9); plt.title("Bilateral");
plt.imshow(pepper_bilateral_cv, cmap="gray"); plt.axis("off")

plt.tight_layout()
plt.show()

```





**Add your analysis and comments here**

## Analysis and Comments

After validating the filtering results using OpenCV's built-in functions, similar behavior was observed compared to the manually implemented filters. For the Gaussian noisy house image, the mean filter reduced noise but introduced noticeable blurring, especially around edges and fine structural details. The Gaussian filter provided smoother results while preserving edges slightly better due to its weighted averaging mechanism.

The median filter reduced noise to some extent but was less effective for Gaussian noise compared to the Gaussian filter. The bilateral filter preserved edges more clearly while still reducing noise, maintaining sharper transitions around the roof and window boundaries.

For the pepper image affected by salt-and-pepper noise, the median filter significantly improved the image quality by removing most of the impulse noise while maintaining object contours. The

bilateral filter preserved edges but did not remove the impulse noise as effectively as the median filter.

Overall, the OpenCV validation confirmed that the behavior of the filters was consistent with theoretical expectations: linear filters smoothed noise but blurred edges, while non-linear filters—particularly the median filter—were more effective for impulse noise and better at preserving structural details.

### T3.3: Perform qualitative evaluation for image denoising

In this task, you will assess the effectiveness of different filtering techniques in removing noise from images through visual inspection.

**Use all the given images for this evaluation to ensure a comprehensive analysis across different image types.**

Steps to follow:

(Steps 1, 2, and 3 have already been completed in previous tasks.)

1. Start with a clean image (ground truth).
2. Add artificial noise to the image to simulate real-world noise conditions.
3. Apply each filtering method one by one to reduce or remove the noise.
4. Display the images in the following order : Original image --> Noisy image --> Filtered image.
5. Visually inspect the filtered images and compare them to the original (ground truth).
6. Observe how well each filter removes noise while preserving image quality, including details, textures, and sharpness.
7. For each type of noise, identify the filter that gives the best visual result, balancing noise removal and image clarity.

```
Your code here.
----- Load clean (ground truth) images -----
house = cv2.imread("images/house.png", cv2.IMREAD_GRAYSCALE)
pepper = cv2.imread("images/pepper.bmp", cv2.IMREAD_GRAYSCALE)

if house is None or pepper is None:
 raise FileNotFoundError("Check paths: images/house.png and
images/pepper.bmp")

----- Load noisy images (saved from Task 2)

house_noisy = cv2.imread("images/house_gaussian_noise.png",
cv2.IMREAD_GRAYSCALE)
pepper_noisy = cv2.imread("images/pepper_salt_pepper_noise.png",
cv2.IMREAD_GRAYSCALE)

if house_noisy is None or pepper_noisy is None:
 raise FileNotFoundError("Noisy images not found. Check if you
saved them in images/ folder.")
```

```

----- Apply filters (OpenCV built-in) -----
Linear filters
house_mean = cv2.blur(house_noisy, (3, 3))
house_gauss = cv2.GaussianBlur(house_noisy, (3, 3), 0)

pepper_mean = cv2.blur(pepper_noisy, (3, 3))
pepper_gauss = cv2.GaussianBlur(pepper_noisy, (3, 3), 0)

Non-linear filters
house_median = cv2.medianBlur(house_noisy, 3)
house_bilat = cv2.bilateralFilter(house_noisy, d=3, sigmaColor=50,
sigmaSpace=50)

pepper_median = cv2.medianBlur(pepper_noisy, 3)
pepper_bilat = cv2.bilateralFilter(pepper_noisy, d=3, sigmaColor=50,
sigmaSpace=50)


----- Display in required order: Original -> Noisy ->
Filtered -----
House (Gaussian noise)
plt.figure(figsize=(14, 8))

plt.subplot(2, 6, 1); plt.title("House (Original)"); plt.imshow(house,
cmap="gray"); plt.axis("off")
plt.subplot(2, 6, 2); plt.title("House (Noisy)");
plt.imshow(house_noisy, cmap="gray"); plt.axis("off")
plt.subplot(2, 6, 3); plt.title("Mean"); plt.imshow(house_mean,
cmap="gray"); plt.axis("off")
plt.subplot(2, 6, 4); plt.title("Gaussian"); plt.imshow(house_gauss,
cmap="gray"); plt.axis("off")
plt.subplot(2, 6, 5); plt.title("Median"); plt.imshow(house_median,
cmap="gray"); plt.axis("off")
plt.subplot(2, 6, 6); plt.title("Bilateral"); plt.imshow(house_bilat,
cmap="gray"); plt.axis("off")

Pepper (Salt & Pepper noise)
plt.subplot(2, 6, 7); plt.title("Pepper (Original)");
plt.imshow(pepper, cmap="gray"); plt.axis("off")
plt.subplot(2, 6, 8); plt.title("Pepper (Noisy)");
plt.imshow(pepper_noisy, cmap="gray"); plt.axis("off")
plt.subplot(2, 6, 9); plt.title("Mean"); plt.imshow(pepper_mean,
cmap="gray"); plt.axis("off")
plt.subplot(2, 6, 10); plt.title("Gaussian"); plt.imshow(pepper_gauss,
cmap="gray"); plt.axis("off")
plt.subplot(2, 6, 11); plt.title("Median"); plt.imshow(pepper_median,
cmap="gray"); plt.axis("off")
plt.subplot(2, 6, 12); plt.title("Bilateral"); plt.imshow(pepper_bilat,
cmap="gray"); plt.axis("off")

```

```
plt.tight_layout()
plt.show()
```



**Add your analysis and comments here**

## Analysis and Comments

The qualitative evaluation showed clear differences in how each filter handled the two types of noise. For the house image corrupted with Gaussian noise, the mean filter reduced noise but introduced noticeable blurring, especially along roof edges and window boundaries. The Gaussian filter achieved smoother noise reduction while preserving edges slightly better due to its weighted averaging mechanism. The median filter reduced some noise but was less effective for Gaussian noise compared to the Gaussian filter. The bilateral filter maintained sharper edges while still reducing noise, resulting in a better balance between smoothness and structural preservation.

For the pepper image affected by salt-and-pepper noise, the median filter produced the most visually satisfying result. Most impulse noise pixels were removed while object contours and textures remained relatively clear. In contrast, the mean and Gaussian filters blurred the image but did not effectively eliminate the impulse noise. The bilateral filter preserved edges but was less effective than the median filter in removing the extreme black and white pixels.

Overall, the evaluation demonstrated that the effectiveness of a filter depends strongly on the noise type. Gaussian filtering performed better for Gaussian noise, while median filtering was more suitable for salt-and-pepper noise. This confirmed that selecting an appropriate filtering technique is essential for achieving optimal denoising results while preserving image quality.

### T3.4: Perform quantitative evaluation for image denoising

#### Sample Image

In the previous task, you performed a qualitative (subjective) evaluation of the denoising results. Now, you will confirm and quantify our observations using objective evaluation metrics.

## Why Use Evaluation Metrics?

Visual inspection is subjective, it depends on human perception, and different observers might interpret results differently. To make a more reliable assessment, you will use two different quantitative metrics that provide numerical values to measure how close the denoised image is to the original (ground truth).

### 1. Peak Signal-to-Noise Ratio (PSNR):

- Measures the difference between the ground truth image and the denoised image.
- A higher PSNR value indicates better denoising performance, meaning less distortion and noise remaining.
- It is expressed in decibels (dB), with typical values ranging between 20-50 dB, depending on image quality.

### 2. Structural Similarity Index (SSIM):

- Unlike PSNR, which only considers pixel-wise differences, SSIM evaluates the structural similarity between two images.
- It measures how well the contrast, luminance, and texture are preserved in the denoised image.
- SSIM values range from 0 to 1, where 1 indicates a perfect match to the original image.

## What you need to do:

- Use prebuilt functions for PSNR and SSIM from `OpenCV`, `sklearn`, or any other relevant library.
- Compute PSNR and SSIM for each denoised image
- During the evaluation process, use these metrics carefully, especially when the original and filtered images do not have the same size.
- Compare the values across different filters and noise types.
- Identify the filter that provides the highest PSNR and SSIM, confirming the most effective denoising method.

```
import sys
!{sys.executable} -m pip install -U scikit-image
```

Requirement already satisfied: scikit-image in c:\Users\yoursr\AppData\Local\Programs\Python\Python312\Lib\site-packages (0.26.0)  
Requirement already satisfied: numpy>=1.24 in c:\Users\yoursr\AppData\Local\Programs\Python\Python312\Lib\site-packages (from scikit-image) (1.26.4)  
Requirement already satisfied: scipy>=1.11.4 in c:\Users\yoursr\AppData\Local\Programs\Python\Python312\Lib\site-packages (from scikit-image) (1.12.0)  
Requirement already satisfied: networkx>=3.0 in c:\Users\yoursr\AppData\Local\Programs\Python\Python312\Lib\site-packages (from scikit-image) (3.5)

Requirement already satisfied: pillow>=10.1 in c:\Users\yoursr\AppData\Local\Programs\Python\Python312\Lib\site-packages (from scikit-image) (10.2.0)

Requirement already satisfied: imageio!=2.35.0,>=2.33 in c:\Users\yoursr\AppData\Local\Programs\Python\Python312\Lib\site-packages (from scikit-image) (2.37.2)

Requirement already satisfied: tifffile>=2022.8.12 in c:\Users\yoursr\AppData\Local\Programs\Python\Python312\Lib\site-packages (from scikit-image) (2026.2.24)

Requirement already satisfied: packaging>=21 in C:\Users\yoursr\AppData\Roaming\Python\Python312\site-packages (from scikit-image) (25.0)

Requirement already satisfied: lazy-loader>=0.4 in c:\Users\yoursr\AppData\Local\Programs\Python\Python312\Lib\site-packages (from scikit-image) (0.4)

WARNING: Ignoring invalid distribution ~ensorflow-intel (c:\Users\yoursr\AppData\Local\Programs\Python\Python312\Lib\site-packages)

WARNING: Ignoring invalid distribution ~ensorflow-intel (c:\Users\yoursr\AppData\Local\Programs\Python\Python312\Lib\site-packages)

WARNING: Ignoring invalid distribution ~ensorflow-intel (c:\Users\yoursr\AppData\Local\Programs\Python\Python312\Lib\site-packages)

*# T3.4 Quantitative evaluation (PSNR + SSIM) – style similar to the example*

```
import cv2
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import math
```

```
If you get: ModuleNotFoundError: No module named 'skimage'
run once in a separate cell:
!pip install scikit-image
from skimage.metrics import structural_similarity as ssim
```

```
----- PSNR -----
def calculate_psnr(original, filtered):
 # Ensure same size (crop to smallest if needed)
 if original.shape != filtered.shape:
 min_h = min(original.shape[0], filtered.shape[0])
 min_w = min(original.shape[1], filtered.shape[1])
 original = original[:min_h, :min_w]
 filtered = filtered[:min_h, :min_w]

 original = original.astype(np.float32)
 filtered = filtered.astype(np.float32)
```



```

mse = np.mean((original - filtered) ** 2)
if mse == 0:
 return float("inf")

max_pixel = 255.0
psnr = 20 * math.log10(max_pixel / math.sqrt(mse))
return psnr

----- SSIM -----
def calculate_ssim(original, filtered):
 # Ensure same size (crop to smallest if needed)
 if original.shape != filtered.shape:
 min_h = min(original.shape[0], filtered.shape[0])
 min_w = min(original.shape[1], filtered.shape[1])
 original = original[:min_h, :min_w]
 filtered = filtered[:min_h, :min_w]

 return ssim(original, filtered, data_range=255)

----- Load ground truth + noisy images -----
house_gt = cv2.imread("images/house.png", cv2.IMREAD_GRAYSCALE)
pepper_gt = cv2.imread("images/pepper.bmp", cv2.IMREAD_GRAYSCALE)

house_noisy = cv2.imread("images/house_gaussian_noise.png",
cv2.IMREAD_GRAYSCALE)
pepper_noisy = cv2.imread("images/pepper_salt_pepper_noise.png",
cv2.IMREAD_GRAYSCALE)

if house_gt is None or pepper_gt is None:
 raise FileNotFoundError("Ground truth images not found:
images/house.png and/or images/pepper.bmp")
if house_noisy is None or pepper_noisy is None:
 raise FileNotFoundError("Noisy images not found:
images/house_gaussian_noise.png and/or
images/pepper_salt_pepper_noise.png")

----- Apply denoising filters (3x3) -----
House (Gaussian noise)
house_mean = cv2.blur(house_noisy, (3, 3))
house_gaussian = cv2.GaussianBlur(house_noisy, (3, 3), sigmaX=1.0)
house_median = cv2.medianBlur(house_noisy, 3)
house_bilateral = cv2.bilateralFilter(house_noisy, d=3, sigmaColor=50,
sigmaSpace=50)

Pepper (Salt & Pepper noise)

```

```

pepper_mean = cv2.blur(pepper_noisy, (3, 3))
pepper_gaussian = cv2.GaussianBlur(pepper_noisy, (3, 3), sigmaX=1.0)
pepper_median = cv2.medianBlur(pepper_noisy, 3)
pepper_bilateral = cv2.bilateralFilter(pepper_noisy, d=3,
sigmaColor=50, sigmaSpace=50)

----- Print results in a table-like format

print("=== QUANTITATIVE EVALUATION: Gaussian Noise (House Image) ===\n")
print(f"{'Filter Type':<20} {'PSNR (dB)':<15} {'SSIM (0-1)':<15}\n{'Noise Type'}")
print("-" * 65)

Baseline (Noisy)
psnr_house_noisy = calculate_psnr(house_gt, house_noisy)
ssim_house_noisy = calculate_ssim(house_gt, house_noisy)
print(f"{'Noisy Image':<20} {psnr_house_noisy:<15.2f}\n{ssim_house_noisy:<15.4f} {'Gaussian'}")

Filters
house_filters = {
 "Mean Filter": house_mean,
 "Gaussian Filter": house_gaussian,
 "Median Filter": house_median,
 "Bilateral Filter": house_bilateral
}

psnr_house_vals, ssim_house_vals = [], []
for fname, fimg in house_filters.items():
 p = calculate_psnr(house_gt, fimg)
 s = calculate_ssim(house_gt, fimg)
 psnr_house_vals.append(p)
 ssim_house_vals.append(s)
 print(f"{'fname':<20} {p:<15.2f} {s:<15.4f} {'Gaussian'}")

print("\n" + "=" * 65 + "\n")

print("=== QUANTITATIVE EVALUATION: Salt & Pepper Noise (Pepper Image) ===\n")
print(f"{'Filter Type':<20} {'PSNR (dB)':<15} {'SSIM (0-1)':<15}\n{'Noise Type'}")
print("-" * 65)

Baseline (Noisy)
psnr_pepper_noisy = calculate_psnr(pepper_gt, pepper_noisy)
ssim_pepper_noisy = calculate_ssim(pepper_gt, pepper_noisy)
print(f"{'Noisy Image':<20} {psnr_pepper_noisy:<15.2f}\n{ssim_pepper_noisy:<15.4f} {'Salt & Pepper'}")

```



```

pepper_filters = {
 "Mean Filter": pepper_mean,
 "Gaussian Filter": pepper_gaussian,
 "Median Filter": pepper_median,
 "Bilateral Filter": pepper_bilateral
}

psnr_pepper_vals, ssim_pepper_vals = [], []
for fname, fimg in pepper_filters.items():
 p = calculate_psnr(pepper_gt, fimg)
 s = calculate_ssim(pepper_gt, fimg)
 psnr_pepper_vals.append(p)
 ssim_pepper_vals.append(s)
 print(f"{fname:<20} {p:<15.2f} {s:<15.4f} {'Salt & Pepper'}")

print("\n" + "=" * 65)

----- Optional: Put results into a DataFrame

rows = []
for (img_name, noise_type, gt, noisy, filt_dict) in [
 ("House", "Gaussian", house_gt, house_noisy, house_filters),
 ("Pepper", "Salt & Pepper", pepper_gt, pepper_noisy,
 pepper_filters)
]:
 # baseline
 rows.append([img_name, "Noisy Image", noise_type,
 calculate_psnr(gt, noisy), calculate_ssim(gt, noisy)])
 # filters
 for fname, fimg in filt_dict.items():
 rows.append([img_name, fname, noise_type, calculate_psnr(gt,
 fimg), calculate_ssim(gt, fimg)])

df = pd.DataFrame(rows, columns=["Image", "Filter", "Noise", "PSNR
(dB)", "SSIM"])
display(df)

----- Best filters -----
filter_names = ["Mean", "Gaussian", "Median", "Bilateral"]

best_house_psnr_idx = int(np.argmax(psnr_house_vals))
best_house_ssim_idx = int(np.argmax(ssim_house_vals))

best_pepper_psnr_idx = int(np.argmax(psnr_pepper_vals))
best_pepper_ssim_idx = int(np.argmax(ssim_pepper_vals))

print("\n=== BEST FILTERS BY METRICS ===")

```

```

print("\nFor Gaussian Noise (House Image):")
print(f"Best PSNR: {filter_names[best_house_psnr_idx]} Filter ({psnr_house_vals[best_house_psnr_idx]:.2f} dB)")
print(f"Best SSIM: {filter_names[best_house_ssim_idx]} Filter ({ssim_house_vals[best_house_ssim_idx]:.4f})")

print("\nFor Salt & Pepper Noise (Pepper Image):")
print(f"Best PSNR: {filter_names[best_pepper_psnr_idx]} Filter ({psnr_pepper_vals[best_pepper_psnr_idx]:.2f} dB)")
print(f"Best SSIM: {filter_names[best_pepper_ssim_idx]} Filter ({ssim_pepper_vals[best_pepper_ssim_idx]:.4f})")

----- Optional: simple plots (PSNR + SSIM)

filters_plot = ["Noisy", "Mean", "Gaussian", "Median", "Bilateral"]
x = np.arange(len(filters_plot))
width = 0.35

plt.figure(figsize=(14, 5))

PSNR
plt.subplot(1, 2, 1)
psnr_house_plot = [psnr_house_noisy] + psnr_house_vals
psnr_pepper_plot = [psnr_pepper_noisy] + psnr_pepper_vals
plt.bar(x - width/2, psnr_house_plot, width, label="Gaussian Noise (House)")
plt.bar(x + width/2, psnr_pepper_plot, width, label="Salt & Pepper (Pepper)")
plt.xticks(x, filters_plot)
plt.ylabel("PSNR (dB)")
plt.title("PSNR Comparison Across Filters")
plt.legend()
plt.grid(True, alpha=0.3)

SSIM
plt.subplot(1, 2, 2)
ssim_house_plot = [ssim_house_noisy] + ssim_house_vals
ssim_pepper_plot = [ssim_pepper_noisy] + ssim_pepper_vals
plt.bar(x - width/2, ssim_house_plot, width, label="Gaussian Noise (House)")
plt.bar(x + width/2, ssim_pepper_plot, width, label="Salt & Pepper (Pepper)")
plt.xticks(x, filters_plot)
plt.ylabel("SSIM (0-1)")
plt.title("SSIM Comparison Across Filters")
plt.legend()
plt.grid(True, alpha=0.3)

```

```
plt.tight_layout()
plt.show()
```

=== QUANTITATIVE EVALUATION: Gaussian Noise (House Image) ===

Filter Type	PSNR (dB)	SSIM (0-1)	Noise Type
-----			
Noisy Image	20.05	0.2888	Gaussian
Mean Filter	27.63	0.5966	Gaussian
Gaussian Filter	27.70	0.5960	Gaussian
Median Filter	26.45	0.5255	Gaussian
Bilateral Filter	25.28	0.4773	Gaussian

=====

=== QUANTITATIVE EVALUATION: Salt & Pepper Noise (Pepper Image) ===

Filter Type	PSNR (dB)	SSIM (0-1)	Noise Type
-----			
Noisy Image	18.39	0.4265	Salt & Pepper
Mean Filter	25.34	0.6872	Salt & Pepper
Gaussian Filter	25.55	0.6837	Salt & Pepper
Median Filter	32.76	0.9599	Salt & Pepper
Bilateral Filter	18.96	0.4562	Salt & Pepper


=====

	Image	Filter	Noise	PSNR (dB)	SSIM
0	House	Noisy Image	Gaussian	20.048286	0.288760
1	House	Mean Filter	Gaussian	27.633678	0.596643
2	House	Gaussian Filter	Gaussian	27.700649	0.596038
3	House	Median Filter	Gaussian	26.446789	0.525513
4	House	Bilateral Filter	Gaussian	25.283061	0.477317
5	Pepper	Noisy Image	Salt & Pepper	18.387570	0.426546
6	Pepper	Mean Filter	Salt & Pepper	25.338057	0.687208
7	Pepper	Gaussian Filter	Salt & Pepper	25.547866	0.683678
8	Pepper	Median Filter	Salt & Pepper	32.762378	0.959875
9	Pepper	Bilateral Filter	Salt & Pepper	18.961759	0.456151

=== BEST FILTERS BY METRICS ===

For Gaussian Noise (House Image):  
Best PSNR: Gaussian Filter (27.70 dB)  
Best SSIM: Mean Filter (0.5966)

For Salt & Pepper Noise (Pepper Image):  
Best PSNR: Median Filter (32.76 dB)  
Best SSIM: Median Filter (0.9599)



**Add your analysis and comments here**

## Analysis and Comments

From the quantitative evaluation, it was observed that all filtering methods improved the image quality compared to the noisy versions. Both PSNR and SSIM increased after applying the filters, confirming that noise was effectively reduced.

For Gaussian noise (House image), the linear filters performed better overall. The Gaussian and Mean filters provided stronger improvements, indicating that smoothing-based approaches are well suited for noise that is continuously distributed across the image. The Median and Bilateral filters also reduced noise, but their performance was slightly lower in this case.

For Salt & Pepper noise (Pepper image), the Median filter clearly provided the best results. It effectively removed the impulse disturbances while preserving structural details and edges. This confirms that the Median filter is particularly appropriate for impulse-type noise. The Mean and Gaussian filters reduced noise to some extent, but they were less effective in handling isolated extreme pixel values.

Overall, the quantitative metrics confirmed the qualitative observations made previously. The effectiveness of each filter strongly depended on the type of noise applied, demonstrating the importance of selecting a denoising method according to the noise model.

## Task 4: Apply and evaluate spatial domain filterin for edge detection

### T4.1: Implement and compare filtering methods

In this task, you will implement and compare different edge detection techniques.

1. Use your 2D convolution function to apply the following edge detection filters:
  - Prewitt filter : Uses two kernels to detect edges in horizontal and vertical directions.

```
----- Load image (grayscale) -----
img = cv2.imread("images/cameraman.jpg", cv2.IMREAD_GRAYSCALE)
if img is None:
 raise FileNotFoundError("Image not found. Check path:
images/cameraman.jpg")

----- Manual VALID convolution (no padding) -----
def conv2d_valid(image, kernel):
 H, W = image.shape
 kH, kW = kernel.shape
 out_H = H - kH + 1
 out_W = W - kW + 1

 output = np.zeros((out_H, out_W), dtype=np.float32)

 # flip kernel for true convolution
 k = np.flipud(np.fliplr(kernel))

 for i in range(out_H):
 for j in range(out_W):
 region = image[i:i+kH, j:j+kW]
 output[i, j] = np.sum(region * k)

 return output

----- Kernels -----
Prewitt
prewitt_x = np.array([[-1, 0, 1],
 [-1, 0, 1],
 [-1, 0, 1]], dtype=np.float32)

prewitt_y = np.array([[-1, -1, -1],
 [0, 0, 0],
 [1, 1, 1]], dtype=np.float32)
```

```

Sobel
sobel_x = np.array([[-1, 0, 1],
 [-2, 0, 2],
 [-1, 0, 1]], dtype=np.float32)

sobel_y = np.array([[-1, -2, -1],
 [0, 0, 0],
 [1, 2, 1]], dtype=np.float32)

Laplacian (4-neighborhood)
laplacian = np.array([[0, -1, 0],
 [-1, 4, -1],
 [0, -1, 0]], dtype=np.float32)

----- Apply edge detectors -----
Prewitt gradients + magnitude
Px = conv2d_valid(img, prewitt_x)
Py = conv2d_valid(img, prewitt_y)
Pmag = np.sqrt(Px**2 + Py**2)

Sobel gradients + magnitude
Sx = conv2d_valid(img, sobel_x)
Sy = conv2d_valid(img, sobel_y)
Smag = np.sqrt(Sx**2 + Sy**2)

Laplacian response (absolute value for visualization)
L = conv2d_valid(img, laplacian)
Labs = np.abs(L)

----- Normalize for display -----
def normalize_to_u8(x):
 x = x - x.min()
 if x.max() != 0:
 x = x / x.max()
 return (x * 255).astype(np.uint8)

Px_u8, Py_u8, Pmag_u8 = normalize_to_u8(np.abs(Px)),
normalize_to_u8(np.abs(Py)), normalize_to_u8(Pmag)
Sx_u8, Sy_u8, Smag_u8 = normalize_to_u8(np.abs(Sx)),
normalize_to_u8(np.abs(Sy)), normalize_to_u8(Smag)
L_u8 = normalize_to_u8(Labs)

----- Display -----
plt.figure(figsize=(14, 10))

plt.subplot(3, 3, 1); plt.title("Original"); plt.imshow(img,
cmap="gray"); plt.axis("off")

plt.subplot(3, 3, 2); plt.title("Prewitt X"); plt.imshow(Px_u8,
cmap="gray"); plt.axis("off")

```




```
plt.subplot(3, 3, 3); plt.title("Prewitt Y"); plt.imshow(Py_u8,
cmap="gray"); plt.axis("off")
plt.subplot(3, 3, 5); plt.title("Prewitt Magnitude");
plt.imshow(Pmag_u8, cmap="gray"); plt.axis("off")

plt.subplot(3, 3, 6); plt.title("Sobel X"); plt.imshow(Sx_u8,
cmap="gray"); plt.axis("off")
plt.subplot(3, 3, 7); plt.title("Sobel Y"); plt.imshow(Sy_u8,
cmap="gray"); plt.axis("off")
plt.subplot(3, 3, 8); plt.title("Sobel Magnitude");
plt.imshow(Smag_u8, cmap="gray"); plt.axis("off")

plt.subplot(3, 3, 9); plt.title("Laplacian"); plt.imshow(L_u8,
cmap="gray"); plt.axis("off")

plt.tight_layout()
plt.show()

Shapes (valid convolution reduces size)
print("Original shape:", img.shape)
print("Prewitt magnitude shape:", Pmag.shape)
print("Sobel magnitude shape:", Smag.shape)
print("Laplacian shape:", L.shape)
```



Original shape: (473, 473)  
Prewitt magnitude shape: (471, 471)  
Sobel magnitude shape: (471, 471)  
Laplacian shape: (471, 471)

**Add your analysis and comments here**

## Analysis and Comments

The edge detection results showed clear differences between the applied operators. The Prewitt and Sobel filters successfully detected horizontal and vertical edges by computing gradients in both directions and combining them using the gradient magnitude. The main structural contours of the cameraman and the tripod were clearly visible in both cases.

It was observed that the Sobel operator produced slightly stronger and more defined edges compared to Prewitt. This is likely due to the higher weighting of central pixels in the Sobel kernels, which improves edge sensitivity while maintaining stability.



The Laplacian operator detected edges without considering direction, highlighting intensity changes in all orientations. However, it also appeared more sensitive to noise, resulting in a more textured and less clean edge map compared to the gradient-based methods.

Additionally, the output image sizes were smaller than the original image, confirming that valid convolution (without padding) was used. Overall, Sobel provided the most visually balanced edge detection, while Laplacian emphasized fine details but introduced more noise-like artifacts.

## T4.2: Validate Edge Detection with OpenCV

In this task, you will apply the same edge detection filters using OpenCV's built-in functions and compare them with your manual implementation.

- Use OpenCV's built-in functions for edge detection:
  - `cv2.filter2D()`: For Prewitt, Sobel and Laplacian filters using convolution operation.
  - `cv2.Canny()`: Another method for edge detection using the Canny method.
- Compare the results of manual implementation with OpenCV's functions.

```
Your code here.
----- Load image -----
img = cv2.imread("images/cameraman.jpg", cv2.IMREAD_GRAYSCALE)
if img is None:
 raise FileNotFoundError("Image not found. Check path:
images/cameraman.jpg")

----- Manual VALID convolution (same as T4.1) -----
def conv2d_valid(image, kernel):
 H, W = image.shape
 kH, kW = kernel.shape
 out_H = H - kH + 1
 out_W = W - kW + 1
 out = np.zeros((out_H, out_W), dtype=np.float32)

 k = np.flipud(np.fliplr(kernel))
 for i in range(out_H):
 for j in range(out_W):
 region = image[i:i+kH, j:j+kW]
 out[i, j] = np.sum(region * k)
 return out

def normalize_to_u8(x):
 x = x.astype(np.float32)
 x = x - x.min()
 if x.max() != 0:
 x = x / x.max()
 return (x * 255).astype(np.uint8)

----- Kernels -----
prewitt_x = np.array([[-1, 0, 1],
 [-1, 0, 1],
```

```

 [-1, 0, 1]], dtype=np.float32)
prewitt_y = np.array([[-1, -1, -1],
 [0, 0, 0],
 [1, 1, 1]], dtype=np.float32)

sobel_x = np.array([[-1, 0, 1],
 [-2, 0, 2],
 [-1, 0, 1]], dtype=np.float32)
sobel_y = np.array([[-1, -2, -1],
 [0, 0, 0],
 [1, 2, 1]], dtype=np.float32)

laplacian_k = np.array([[0, -1, 0],
 [-1, 4, -1],
 [0, -1, 0]], dtype=np.float32)

----- Manual edges -----
Px_m = conv2d_valid(img, prewitt_x)
Py_m = conv2d_valid(img, prewitt_y)
Pmag_m = np.sqrt(Px_m**2 + Py_m**2)

Sx_m = conv2d_valid(img, sobel_x)
Sy_m = conv2d_valid(img, sobel_y)
Smag_m = np.sqrt(Sx_m**2 + Sy_m**2)

L_m = conv2d_valid(img, laplacian_k)
Labs_m = np.abs(L_m)

----- OpenCV edges -----
filter2D returns same size as input, so we crop to "valid" region
for comparison
Px_cv_full = cv2.filter2D(img, cv2.CV_32F, prewitt_x)
Py_cv_full = cv2.filter2D(img, cv2.CV_32F, prewitt_y)
Sx_cv_full = cv2.filter2D(img, cv2.CV_32F, sobel_x)
Sy_cv_full = cv2.filter2D(img, cv2.CV_32F, sobel_y)
L_cv_full = cv2.filter2D(img, cv2.CV_32F, laplacian_k)

Crop OpenCV outputs to match manual valid size (remove 1 pixel
border for 3x3)
Px_cv = Px_cv_full[1:-1, 1:-1]
Py_cv = Py_cv_full[1:-1, 1:-1]
Sx_cv = Sx_cv_full[1:-1, 1:-1]
Sy_cv = Sy_cv_full[1:-1, 1:-1]
L_cv = L_cv_full[1:-1, 1:-1]

Pmag_cv = np.sqrt(Px_cv**2 + Py_cv**2)
Smag_cv = np.sqrt(Sx_cv**2 + Sy_cv**2)
Labs_cv = np.abs(L_cv)

----- Numeric check (max absolute difference) -----

```

```

print("Max diff Prewitt magnitude:", np.max(np.abs(Pmag_m - Pmag_cv)))
print("Max diff Sobel magnitude:", np.max(np.abs(Smag_m - Smag_cv)))
print("Max diff Laplacian:", np.max(np.abs(Labs_m - Labs_cv)))

----- Canny (OpenCV) -----
Canny works differently (not just a single convolution), so it's
shown separately
canny = cv2.Canny(img, 100, 200)

----- Display comparison -----
plt.figure(figsize=(14, 10))

plt.subplot(3, 3, 1); plt.title("Original"); plt.imshow(img,
cmap="gray"); plt.axis("off")

plt.subplot(3, 3, 2); plt.title("Manual Prewitt Mag");
plt.imshow(normalize_to_u8(Pmag_m), cmap="gray"); plt.axis("off")
plt.subplot(3, 3, 3); plt.title("OpenCV Prewitt Mag");
plt.imshow(normalize_to_u8(Pmag_cv), cmap="gray"); plt.axis("off")

plt.subplot(3, 3, 4); plt.title("Manual Sobel Mag");
plt.imshow(normalize_to_u8(Smag_m), cmap="gray"); plt.axis("off")
plt.subplot(3, 3, 5); plt.title("OpenCV Sobel Mag");
plt.imshow(normalize_to_u8(Smag_cv), cmap="gray"); plt.axis("off")


plt.subplot(3, 3, 6); plt.title("Manual Laplacian");
plt.imshow(normalize_to_u8(Labs_m), cmap="gray"); plt.axis("off")
plt.subplot(3, 3, 7); plt.title("OpenCV Laplacian");
plt.imshow(normalize_to_u8(Labs_cv), cmap="gray"); plt.axis("off")

plt.subplot(3, 3, 8); plt.title("OpenCV Canny"); plt.imshow(canny,
cmap="gray"); plt.axis("off")

plt.tight_layout()
plt.show()

Max diff Prewitt magnitude: 0.0
Max diff Sobel magnitude: 0.0
Max diff Laplacian: 0.0

```



**Add your analysis and comments here**

## Analysis and Comments

The OpenCV results supported the correctness of the manual implementation, and the comparison showed how different edge detectors emphasize edges with different levels of sharpness, noise sensitivity, and structural clarity.

## Task 5: Experiment with other filter size

Now, **re-perform Task 3 and Task 4** using different kernel sizes to observe how the results change.

- Instead of using a **3×3 kernel**, try **5×5**, **7×7**, and **9×9** kernels.
- Use different values of noise density to see **whether** it is effectively filtered across different kernel sizes.
- To avoid unnecessary steps, apply only the most suitable filter for each type of noise based on your findings.
- Compare the results and analyze:

- How does increasing the kernel size affect noise reduction?
- How does it impact edge preservation?

## Task 5: Effect of Kernel Size on Denoising

In this task, the most suitable filter for each noise type will be applied using different kernel sizes (3×3, 5×5, 7×7, 9×9).

We will observe how increasing the kernel size affects noise reduction and edge preservation.

### Experiment 1: Gaussian Noise (Gaussian Filter)

The Gaussian filter will be applied using increasing kernel sizes.

The effect on smoothing strength and edge clarity will be observed.

```
import cv2
import numpy as np
import matplotlib.pyplot as plt
import math

Load ground truth images
house = cv2.imread("images/house.png", cv2.IMREAD_GRAYSCALE)
pepper = cv2.imread("images/pepper.bmp", cv2.IMREAD_GRAYSCALE)

Load noisy images
house_noisy = cv2.imread("images/house_gaussian_noise.png",
cv2.IMREAD_GRAYSCALE)
pepper_noisy = cv2.imread("images/pepper_salt_pepper_noise.png",
cv2.IMREAD_GRAYSCALE)


if house is None or pepper is None or house_noisy is None or
pepper_noisy is None:
 raise FileNotFoundError("Check that all images exist in the
images/ folder.")
kernel_sizes = [3, 5, 7, 9]

plt.figure(figsize=(12,8))
plt.subplot(2,3,1)
plt.title("Noisy Image")
plt.imshow(house_noisy, cmap="gray")
plt.axis("off")

for i, k in enumerate(kernel_sizes):
 filtered = cv2.GaussianBlur(house_noisy, (k,k), 0)

 plt.subplot(2,3,i+2)
 plt.title(f"Gaussian {k}x{k}")
 plt.imshow(filtered, cmap="gray")
 plt.axis("off")
```

```
plt.tight_layout()
plt.show()
```



## Analysis and Comments

As the kernel size increased, stronger smoothing was observed. Noise was reduced more effectively with larger kernels; however, edges became progressively blurred.

Smaller kernels preserved structural details better, while larger kernels removed more noise at the cost of sharpness.

This demonstrates the trade-off between noise reduction and edge preservation.

## Experiment 2: Salt & Pepper Noise (Median Filter)


The Median filter will be applied using increasing kernel sizes. The impact on impulse noise removal and edge clarity will be examined.

```
plt.figure(figsize=(12,8))
plt.subplot(2,3,1)
plt.title("Noisy Image")
plt.imshow(pepper_noisy, cmap="gray")
plt.axis("off")
```

```
for i, k in enumerate(kernel_sizes):
 filtered = cv2.medianBlur(pepper_noisy, k)

 plt.subplot(2,3,i+2)
 plt.title(f"Median {k}x{k}")
 plt.imshow(filtered, cmap="gray")
 plt.axis("off")

plt.tight_layout()
plt.show()
```



## Analysis and Comments

From the results, it was observed that the Median filter progressively reduced the Salt & Pepper noise as the kernel size increased. With the 3×3 kernel, most impulse noise was removed while preserving important structural details of the peppers. The edges and textures remained relatively clear.

When the kernel size increased to 5×5, noise reduction became stronger, and the image appeared smoother. However, slight blurring of fine details was noticeable.

For larger kernel sizes such as 7×7 and 9×9, the image became significantly smoother, but edge sharpness and texture details were reduced. Although noise was almost completely removed, the loss of fine structures became more evident.

Overall, increasing the kernel size improved impulse noise removal, but it also led to reduced edge preservation. A moderate kernel size provided a better balance between effective denoising and maintaining image clarity.

## Final Analysis and Comparison

To better summarize the impact of kernel size on denoising performance, the observations are presented below:

Noise Type	Filter Used	Small Kernel (3×3)	Medium Kernel (5×5 / 7×7)	Large Kernel (9×9)
Gaussian Noise	Gaussian Filter	Moderate smoothing, good detail preservation	Stronger smoothing, slight edge softening	Strong smoothing, noticeable blur
Salt & Pepper	Median Filter	Effective impulse removal, edges preserved	Very effective denoising, slight loss of texture	Very smooth image, reduced sharpness

From the experiments, increasing the kernel size consistently improved noise removal in both Gaussian and Salt & Pepper cases. However, larger kernels also introduced stronger blurring effects and reduced fine structural details.

For Gaussian noise, moderate kernel sizes provided the best balance between smoothing and preserving image clarity. Very large kernels removed more noise but significantly softened edges.

For Salt & Pepper noise, the Median filter effectively eliminated impulse disturbances even with smaller kernels. While larger kernels further reduced residual noise, they also reduced edge sharpness and texture details.

## Conclusion

Kernel size plays a critical role in controlling the trade-off between denoising strength and edge preservation. Smaller kernels maintain structural details but may leave some noise, whereas larger kernels enhance smoothing at the expense of sharpness. Therefore, selecting an appropriate kernel size depends on the noise characteristics and the level of detail that must be preserved in the image.